

Performance Test

Hive/Spark (Parquet) vs. Azure Databricks Delta

Compare performance of Hive/Spark tables
(with underlying Parquet file format)
with
Azure Databricks Delta tables

Prepared by

DM Jumpstart Engineering Team
(askdmjfordmtools@microsoft.com)

Disclaimer

The High-Level Architecture, Migration Dispositions and guidelines in this document is developed in consultation and collaboration with Microsoft Corporation technical architects. Because Microsoft must respond to changing market conditions, this document should not be interpreted as an invitation to contract or a commitment on the part of Microsoft.

Microsoft has provided generic high-level guidance in this document with the understanding that MICROSOFT MAKES NO WARRANTIES, EXPRESS OR IMPLIED, WITH RESPECT TO THE INFORMATION CONTAINED HEREIN.

This document is provided “as-is”. Information and views expressed in this document, including URL and other Internet Web site references, may change without notice.

Some examples depicted herein are provided for illustration only and are fictitious. No real association or connection is intended or should be inferred.

This document does not provide you with any legal rights to any intellectual property in any Microsoft product. You may copy and use this document for your internal, reference purposes.

© 2019 Microsoft. All rights reserved.

Note: The detail provided in this document has been harvested as part of a customer engagement sponsored through the [Azure Data Services Jumpstart Program](#).

What is Databricks Delta?

Delta is an optimized Spark table that brings data reliability and performance optimizations to cloud data lakes. Making your data lake faster and more reliable – accelerating innovation.

Databricks Runtime *is the most optimized version of...*



A Delta Table *is the most optimized version of...* A Spark/Hive Table

<https://docs.azuredatabricks.net/delta/delta-intro.html>

Databricks Delta - Unified Data Management

Delta is an optimized Spark table, **unifying both streaming and batch data**, that stores data in the open-source Parquet format in blob storage.

- **Improves performance** by managing metadata about your table
 - Automated file management to make engineering simple (compaction)
 - Creates and maintains indexes to speed up queries 10x – 100x
- **Improves data reliability & simplifies pipelines** by adding data management capabilities
 - Transactional guarantees (ACID)
 - Schema enforcement ensures data integrity
 - UPSERTs/Deletes allow for changing data

Performance Test

- Cluster Size – 8 Nodes
 - Each node is Standard_DS3_v2 – with 4 cores, 14 GB
 - No autoscaling for the cluster
- Size of parquet files – ~32 GB (uncompressed ~340 GB)
- Statistics – Rows: 635,586,654 – Columns: 84
- Databricks Runtime – 4.3 beta (includes Apache Spark 2.3.1, Scala 2.11)

Please note, although this data volume is not huge enough to represent the big data scenario however it will give an idea of the performance differences and how you can carry out this exercise in your environment with your big data.

Performance Test

Hive External Table – Parquet

```
1 %sql
2 SELECT BusDt, SUM(GrossSalesLclAmt), SUM(NetRevLclAmt)
3 FROM EDW.pos_hive_big
4 GROUP BY BusDt
5 ORDER BY BusDt
```

▶ (1) Spark Jobs

BusDt	sum(GrossSalesLclAmt)
2018-07-01	Sanitized
2018-07-02	
2018-07-03	
2018-07-04	
2018-07-05	
2018-07-06	
2018-07-07	
2018-07-08	
2018-07-09	



Command took 1.36 minutes -- by i m at 8/9/2018 9:15:35 PM on arali_hive2

Delta Table

```
1 %sql
2 SELECT BusDt, SUM(GrossSalesLclAmt), SUM(NetRevLclAmt)
3 FROM EDW.pos_delta_big
4 GROUP BY BusDt
5 ORDER BY BusDt
```

▶ (4) Spark Jobs

BusDt	sum(GrossSalesLclAmt)
2018-07-01	Sanitized
2018-07-02	
2018-07-03	
2018-07-04	
2018-07-05	
2018-07-06	
2018-07-07	
2018-07-08	
2018-07-09	



Command took 1.19 minutes -- by i m at 8/9/2018 9:17:04 PM on arali_hive2

Please note, the given time difference between Hive/Spark table (with underlying Parquet) and Delta table doesn't look huge, however it would be much better with bigger dataset, especially when you leverage delta optimization (discussed later).

Performance Test

Hive External Table – Parquet

```
1 %sql
2 SELECT BusDt, SUM(GrossSalesLclAmt), SUM(NetRevLclAmt)
3 FROM EDW.pos_hive_big
4 WHERE BusDt BETWEEN '2018-07-10' AND '2018-07-25'
5 GROUP BY BusDt
6 ORDER BY BusDt
```

▶ (1) Spark Jobs

BusDt	sum(GrossSalesLclAmt)
2018-07-10	
2018-07-11	
2018-07-12	
2018-07-13	
2018-07-14	
2018-07-15	
2018-07-16	
2018-07-17	
2018-07-18	

Command took 32.66 seconds -- by a :om at 8/9/2018 9:26:14 PM on arali_hive2

Delta Table

```
1 %sql
2 SELECT BusDt, SUM(GrossSalesLclAmt), SUM(NetRevLclAmt)
3 FROM EDW.pos_delta_big
4 WHERE BusDt BETWEEN '2018-07-10' AND '2018-07-25'
5 GROUP BY BusDt
6 ORDER BY BusDt
```

▶ (2) Spark Jobs

BusDt	sum(GrossSalesLclAmt)
2018-07-10	
2018-07-11	
2018-07-12	
2018-07-13	
2018-07-14	
2018-07-15	
2018-07-16	
2018-07-17	
2018-07-18	

Command took 31.32 seconds -- by a :om at 8/9/2018 9:27:31 PM on arali_hive2

Please note, the given time difference between Hive/Spark table (with underlying Parquet) and Delta table doesn't look huge, however it would be much better with bigger dataset, especially when you leverage delta optimization (discussed later).

Performance Test

Hive External Table – Parquet

```
1 %sql
2 SELECT BusDt, SUM(GrossSalesLclAmt), SUM(NetRevLclAmt)
3 FROM EDW.pos_hive_big
4 WHERE CntryCd = 'CA'
5 GROUP BY BusDt
6 ORDER BY BusDt
```

▶ (1) Spark Jobs

BusDt	sum(GrossSalesLclAmt)
2018-07-01	
2018-07-02	
2018-07-03	
2018-07-04	
2018-07-05	
2018-07-06	
2018-07-07	
2018-07-08	
2018-07-09	

Sanitized

Command took 1.25 minutes -- by a :om at 8/9/2018 10:26:38 PM on arali_hive2

Delta Table

```
1 %sql
2 SELECT BusDt, SUM(GrossSalesLclAmt), SUM(NetRevLclAmt)
3 FROM EDW.pos_delta_big
4 WHERE CntryCd = 'CA'
5 GROUP BY BusDt
6 ORDER BY BusDt
```

▶ (2) Spark Jobs

BusDt	sum(GrossSalesLclAmt)
2018-07-01	
2018-07-02	
2018-07-03	
2018-07-04	
2018-07-05	
2018-07-06	
2018-07-07	
2018-07-08	
2018-07-09	

Sanitized

Command took 1.16 minutes -- by a :om at 8/9/2018 10:29:33 PM on arali_hive2

Please note, the given time difference between Hive/Spark table (with underlying Parquet) and Delta table doesn't look huge, however it would be much better with bigger dataset, especially when you leverage delta optimization (discussed later).

Performance Test – Optimization with Delta

Databricks Delta supports the ability to optimize the layout of data stored in DBFS with these two algorithms:

- Compaction (bin-packing) – Improves the speed of read queries from a table by coalescing small files into larger ones from large number of small files to small number of larger files.
- ZOrdering – Colocates related information in the same set of files. This co-locality is automatically used by Databricks Delta data-skipping algorithms to dramatically reduce the amount of data that needs to be read.

Best Practice - If you have a large amount of data, we suggest that you limit it to new data for optimization, using partition filters (WHERE clause) when possible for ZOrdering.

```
1 %sql
2 OPTIMIZE EDW.pos_delta_big
3 ZORDER BY (CntryCd)
```

▼ (4) Spark Jobs

- ▶ Job 92 [View](#) (Stages: 2/2, 1 skipped)
- ▶ Job 116 [View](#) (Stages: 1/1)
- ▶ Job 124 [View](#) (Stages: 31/31)
- ▶ Job 125 [View](#) (Stages: 4/4, 1 skipped)

path
null

Command took 41.60 minutes -- by [user] at 8/9/2018 10:30:55 PM on arali_hive2

<https://docs.azuredatabricks.net/delta/optimizations.html>

Performance Test

Hive External Table – Parquet

```
1 %sql
2 SELECT BusDt, SUM(GrossSalesLclAmt), SUM(NetRevLclAmt)
3 FROM EDW.pos_hive_big
4 WHERE CntryCd = 'CA'
5 GROUP BY BusDt
6 ORDER BY BusDt
```

▶ (1) Spark Jobs

BusDt	sum(GrossSalesLclAmt)
2018-07-01	Sanitized
2018-07-02	
2018-07-03	
2018-07-04	
2018-07-05	
2018-07-06	
2018-07-07	
2018-07-08	
2018-07-09	



Command took 1.25 minutes -- by ar... on at 8/9/2018 10:26:38 PM on arali_hive2

Delta Table

```
1 %sql
2 SELECT BusDt, SUM(GrossSalesLclAmt), SUM(NetRevLclAmt)
3 FROM EDW.pos_delta_big
4 WHERE CntryCd = 'CA'
5 GROUP BY BusDt
6 ORDER BY BusDt
```

▶ (2) Spark Jobs

BusDt	sum(GrossSalesLclAmt)
2018-07-01	Sanitized
2018-07-02	
2018-07-03	
2018-07-04	
2018-07-05	
2018-07-06	
2018-07-07	
2018-07-08	
2018-07-09	



Command took 18.43 seconds -- by ar... on at 8/9/2018 11:13:16 PM on arali_hive2

Please note, performance of your query would vary here based on amount of data requested by your query. You can find more information here: <https://docs.azuredatabricks.net/delta/optimizations.html>









Benefits of using Databricks Delta

- New data/partition gets reflected automatically to all the clusters – no need for repairing table or deleting/adding partitions to the table
- No need to refresh table metadata in local cache with REFRESH table command
- Optimization for performance and cost
- Transactional consistent –
 - When rebuilding a partition fails, data remains consistent.
 - Multiple users can modify the dataset simultaneously and see the consistent view
- A single delta table can be target for both batch and stream data ingestion – no need to have separate pipelines for streaming and batch

Benefits of using Databricks Delta – Cont'd

- You can write [MERGE statement](#) merge (insert/update) data into Databricks Delta table
- Databricks Delta table also simplifies operational overhead in comparison with Hive External table as explained next

Using Hive External tables –
connected to external metadata
store (Azure SQL Database)

Data Engineering – Worspace/Cluster	Business Group – Workspace/Cluster
<div>1df2.write.format("parquet").partitionBy("BusDt").save("/mnt/workspaces/Personal/arali/pos/parquet_hive_small")</div> <div>▶ (1) Spark Jobs</div> <div>Command took 6.46 minutes -- by : 8/8/2018 3:31:26 PM on arali_hive1</div>	
<div>1%sql 2CREATE TABLE EDW.pos_hive_small 3USING parquet 4OPTIONS (path "/mnt/workspaces/Personal/arali/pos/parquet_hive_small")</div> <div>▶ (1) Spark Jobs</div> <div>OK</div> <div>Command took 3.32 seconds -- by ar :om at 8/8/2018 3:38:53 PM on arali_hive1</div>	
<div>1%sql 2SELECT count(*) from EDW.pos_hive_small</div> <div>▶ (1) Spark Jobs</div> <div>count(1)</div> <div>0</div> <div></div> <div>Command took 1.18 seconds -- by ar :om at 8/9/2018 11:28:27 AM on arali_hive1</div>	<div>1%sql 2SELECT count(*) from EDW.pos_hive_small</div> <div>▶ (1) Spark Jobs</div> <div>count(1)</div> <div>0</div> <div></div> <div>Command took 1.28 seconds -- by a :om at 8/9/2018 11:22:87 AM on arali_hive2</div>
<div>1%sql 2MSCK REPAIR TABLE EDW.pos_hive_small</div> <div>OK</div> <div>Command took 8.85 seconds -- by ar :om at 8/9/2018 11:22:35 AM on arali_hive1</div>	

Data Engineering – Worspace/Cluster

```
1 %sql
2 SELECT count(*) from EDW.pos_hive_small
```

▶ (1) Spark Jobs

count(1)

22701090



Command took 12.88 seconds -- by ar at 8/9/2018 11:23:01 AM on arali_hive1

```
1 df2.write.format("parquet").partitionBy("BusDt").mode("Append").save("/mnt/workspaces/Personal/arali/pos/parquet_hive_small")
```

▶ (1) Spark Jobs

Command took 1.91 minutes -- by /9/2018 11:25:30 AM on arali_hive1

```
1 %sql
2 SELECT count(*) from EDW.pos_hive_small
```

▶ (1) Spark Jobs

count(1)

22701090



Command took 8.22 seconds -- by ar at 8/9/2018 11:29:46 AM on arali_hive1

```
1 %sql
2 MSCK REPAIR TABLE EDW.pos_hive_small
```

OK

Command took 0.96 seconds -- by a at 8/9/2018 11:32:50 AM on arali_hive1

```
1 %sql
2 SELECT count(*) from EDW.pos_hive_small
```

▶ (1) Spark Jobs

count(1)

43312969



Command took 4.42 seconds -- by a at 8/9/2018 11:33:30 AM on arali_hive1

Business Group – Workspace/Cluster

```
1 %sql
2 SELECT count(*) from EDW.pos_hive_small
```

▶ (1) Spark Jobs

count(1)

22701090



Command took 13.02 seconds -- by a at 8/9/2018 11:23:40 AM on arali_hive2

```
1 %sql
2 SELECT count(*) from EDW.pos_hive_small
```

▶ (1) Spark Jobs

count(1)

22701090



Command took 5.08 seconds -- by ar at 8/9/2018 11:30:25 AM on arali_hive2

```
1 %sql
2 SELECT count(*) from EDW.pos_hive_small
```

▶ (1) Spark Jobs

count(1)

43312969



Command took 7.21 seconds -- by a at 8/9/2018 11:34:01 AM on arali_hive2

Data Changes – REFRESH TABLE

```
1 %sql
2 SELECT BusDt, count(*) from EDW.pos_hive_small_wopartition GROUP BY BusDt
```

▶ (1) Spark Jobs

Error in SQL statement: SparkException: Job aborted due to stage failure: Task 0 in stage 35.0 failed 4 times, most recent failure: Lost task 0.3 in stage 35.0 (TID 644, 10.139.64.6, executor 1): java.io.FileNotFoundException: /mnt/workspaces/Personal/arali/pos/parquet_hive_small_wopartition/part-00049-tid-3570897718831618244-b8ec4699-c44b-45b0-8202-6056a4be89ab-670-c000.snappy.parquet

It is possible the underlying files have been updated. You can explicitly invalidate the cache in Spark by running 'REFRESH TABLE tableName' command in SQL or by recreating the Dataset/DataFrame involved.

```
at org.apache.spark.sql.execution.datasources.FileScanRDD$$anon$1.org$apache$spark$sql$execution$datasources$FileScanRDD$$anon$$readFile(FileScanRDD.scala:211)
at org.apache.spark.sql.execution.datasources.FileScanRDD$$anon$1.org$apache$spark$sql$execution$datasources$FileScanRDD$$anon$$createNextIterator(FileScanRDD.scala:380)
at org.apache.spark.sql.execution.datasources.FileScanRDD$$anon$1$$anonfun$prepareNextFile$1.apply(FileScanRDD.scala:298)
at org.apache.spark.sql.execution.datasources.FileScanRDD$$anon$1$$anonfun$prepareNextFile$1.apply(FileScanRDD.scala:294)
```

```
1 %sql
2 REFRESH TABLE EDW.pos_hive_small_wopartition
```

OK

Command took 0.89 seconds -- by arali@arali-hive2.com at 8/9/2018 11:37:56 AM on arali_hive2

Using Databricks Delta –
connected to external metadata
store (Azure SQL Database)

Data Engineering – Workspace/Cluster

```
1 df2.write.format("delta").mode("Append").save("/mnt/workspaces/Personal/arali/pos/parquet_delta")
```

▶ (2) Spark Jobs

Command took 2.17 minutes -- by arali@starbucks.com at 8/8/2018 10:45:43 AM on arali_delta

```
1 %sql
2 SELECT BusDt, SUM(GrossSalesLclAmt), SUM(NetRevLclAmt) FROM pos_delta GROUP BY BusDt
```

▶ (6) Spark Jobs

BusDt	sum(GrossSalesLclAmt)
2018-07-20	Sanitized
2018-07-21	

Command took 11.99 seconds -- by : 8/2018 10:48:33 AM on arali_delta

Business Group – Workspace/Cluster

```
1 %sql
2 SELECT BusDt, SUM(GrossSalesLclAmt), SUM(NetRevLclAmt) FROM pos_delta GROUP BY BusDt
```

▶ (7) Spark Jobs

BusDt	sum(GrossSalesLclAmt)	sum(NetRevLclAmt)
2018-07-20	Sanitized	
2018-07-21		

Command took 18.46 seconds -- m at 8/8/2018 10:49:17 AM on arali_delta


Feedback and suggestions

If you have feedback or suggestions for improving this data migration asset, please contact the Data Migration Jumpstart Team (askdmjfordmtools@microsoft.com). Thanks for your support!

Note: For additional information about migrating various source databases to Azure, see the [Azure Database Migration Guide](#).

Appendix

Data Pipeline – with Spark/Hive/Parquet



ETL Batch (~hours or days)

Compaction Job - Batch (~hours or days)

Queries on Stale Data (Only Partition Pruning for Data Skipping)

Data Pipeline – simplified with Delta

